

# Statistics Guided Exercise

For using statistics with Python, we will be looking at the Pandas library. Pandas itself is built on top of another library, NumPy, and both have their own data structures. In this exercise, we will go over these data structures, and introduce you to Bokeh, which is a visualisation library you will be using in this exercise and the next for graphs and charts.

## Pandas

Pandas is a Python library for data analysis in Python. It provides some useful functions and data structures for the collection and analysis of data. In particular, we will be making use of the **DataFrame** and **Series** classes.

A **DataFrame** object represents data in a series of rows (individual observations of data) and columns (features or variables) within those data. Each of those rows and columns can be extracted, and they then become a **Series**. We will work through an example to illustrate these concepts.

## Importing Data

There are convenience functions to import data, such as **read\_json** and **read\_csv** which, as their names suggest, will import data which is already in a particular format. For this example, we will import data from the MongoDB database we used in the exercise last week.

For this example, we will import the first 1000 documents (*MongoDB stores data records as BSON documents. BSON is a binary representation of JSON documents, though it contains more data types than JSON*), in the UK collection into a Pandas **DataFrame**. Run the cell below

In [1]:

```
# Convention is to import numpy and pandas with abbreviated names
# This means that instead of using pandas.read_csv, you would use pd.read_csv
import numpy as np
import pandas as pd #넘파이의 상위 라이브러리

from bokeh.io import output_notebook, show #파이썬 시각화 응용프로그램 보케
from bokeh.charts import * #차트를 하위 라이브러리로 만들고 보케의 구성요소가 됩니당

# Import PyMongo, so that we can query some data
# 'mongodb://cpduser:M13pV5woDW@mongodb/health_data' is the location of the data we

from pymongo import MongoClient #몽고DB
client = MongoClient('mongodb://cpduser:M13pV5woDW@mongodb/health_data')
db = client.health_data

cursor = db.uk.find({}).limit(1000) #할당을 할 때 결과에 제한(1000개 넘지 말것) 을 걸었음
#빈 딕셔너리에

# Unfortunately, Pandas does not support PyMongo objects for import, so we need to c
listy = list(cursor) #변수 커서의 유형을 리스트로 설정, 판다스를 써서 쓸 수 있게 하고 있음

# Create a Pandas DataFrame with the list object as a parameter
first_1000 = pd.DataFrame(listy) #리스트인 리스트를 데이터 프레임 형식으로 바꿈
print(first_1000)
```

Line3 \	AddressLine1	AddressLine2	Address
0	NaN	16a Adelphi Street	Pr
eston			
1	NaN	24 Fylde Road	Pr
eston			
2	NaN	119a Church Street	Pr
eston			
3	NaN	30-34 Holmrook Road	Pr
eston			
4	NaN	NaN	
NaN			
5	NaN	109-113 Avenham Lane	Pr
eston			
6	NaN	NaN	
NaN			
7	NaN	515 Blackpool Road	Pr
eston			
8	NaN	27 Meadow Street	Pr
---			

Now we have our imported data in a **DataFrame object**. Like any other Python object, it has a collection of attributes and methods which we can use. We will go over some here, but see [the documentation](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.html) (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.html>) for a full list. We'll start by seeing what the data looks like by calling the `head()` function on the data:

In [29]:

```
# Filtering the data in the DataFrame to only return rows where RatingValue < 3
first_1000[first_1000['RatingValue'] < 3]
#결과 값이 레이팅벨류가 0, 1, 2 인 것만 출력
```

292	NaN	49 Blackpool Road	Preston	NaN	Don Mario	Takeaway/sandwich shop
327	NaN	33-35 Moor Lane	Preston	NaN	Far East Oriental Foodstore	Retailers - other

		303 Blackpool			Fredy's at	
--	--	---------------	--	--	------------	--

In [5]:

```
# We can also create a DataFrame which has only some columns
three_columns = first_1000[['RatingValue', 'FHRSID', 'PostCode']]
#열이 있는 데이터 프레임 생성

print('DataFrame with COLUMNS only:\n\n', three_columns.head())
#헤드는 변수 t_c의 함수가 되며 해당 변수에 적용됨
#--> 헤드의 의미 : 처음 나오는 5개의 레코드를 보여라
# Or some rows
# Print the DataFrame with rows 950 to 960, could be any number of rows
print("\nDataFrame with a selection of ROWS:")
first_1000[950:960]
#행 950~960 출력
```

DataFrame with COLUMNS only:

	RatingValue	FHRSID	PostCode
0	3.0	90105	PR1 7BE
1	5.0	479725	PR1 7BY
2	5.0	479676	PR1 3BT
3	2.0	135423	PR1 6SR
4	5.0	759083	NaN

DataFrame with a selection of ROWS:

Out[5]:

	AddressLine1	AddressLine2	AddressLine3	AddressLine4	BusinessName	Business
950	NaN	Sherwood Way	Preston	NaN	Sherwood Court	Hospitals Premises
951	NaN	Sherwood Way	Preston	NaN	Sherwood Lodge Res. Care Home	Hospitals Premises
952	NaN	14 Avenham Road	Preston	NaN	Shining Stars Nursery	Hospitals Premises
953	NaN	3 Fylde Road	Preston	NaN	Ships & Giggles	Pub/bar/
954	NaN	171 New Hall Lane	Preston	NaN	Shop Rite	Retailers

	AddressLine1	AddressLine2	AddressLine3	AddressLine4	BusinessName	BusinessType
955	NaN	219 Ribbleton Lane	Preston	NaN	Sicilia Uno	Takeawa
956	NaN	NaN	NaN	NaN	Simply Country & Co	Other ca
957	NaN	NaN	NaN	NaN	Simply Delish	Other ca
958	NaN	NaN	NaN	NaN	Sion House	Hospital/ Premises
959	NaN	Blackpool Road	Preston	NaN	Sir Tom Finney Community High School	School/c

10 rows × 29 columns

Using the existing `first_1000` DataFrame, try and create a dataset which outputs the columns `FHRSID`, `PostCode`, `LocalAuthorityName`, with any establishment where `RatingValue < 3`

In [6]:

```
# YOUR CODE HERE
three_columns = first_1000[first_1000['RatingValue']<3][['FHRSID', 'PostCode', 'LocalAuthorityName']]
#뒤에 빨간색 글씨 기준을 적용해야 되기 때문에 first_1000을 다시 입력함
#이제 여기 뒤에 나올 FHRSID, PostCode, LocalAuthorityName 적어주세요
print('DataFrame with COLUMNS only:\n\n', three_columns)
```

DataFrame with COLUMNS only:

	FHRSID	PostCode	LocalAuthorityName
3	135423	PR1 6SR	Preston
19	479641	PR1 2XQ	Preston
64	335295	PR1 3BT	Preston
85	448032	PR1 4DX	Preston
108	454884	PR1 1PX	Preston
120	51647	PR1 6XH	Preston
127	448082	PR1 2XQ	Preston
180	369063	PR2 6NH	Preston
189	69893	PR1 8JD	Preston
206	86973	PR1 1PX	Preston
230	637915	PR1 5LD	Preston
265	335352	PR1 8RQ	Preston
272	121600	PR1 1TS	Preston
273	133196	PR2 2DX	Preston
276	335380	PR1 7BE	Preston
292	335301	PR2 6BU	Preston
327	479591	PR1 7AT	Preston
350	201373	PR2 2DU	Preston
398	629593	PR1 5NU	Preston
426	479519	PR1 4SS	Preston
432	115804	PR1 5HH	Preston
459	479764	PR1 2AR	Preston
469	479921	PR1 4ST	Preston
474	768229	PR2 1AU	Preston
518	137209	PR1 4SU	Preston
561	850075	PR1 7JN	Preston
575	479767	PR2 1XN	Preston
632	201370	PR1 5EA	Preston
648	448112	PR1 7JS	Preston
654	70877	PR1 5QS	Preston
663	97856	PR1 8HJ	Preston
665	708327	PR1 7RA	Preston
687	335282	PR1 3BQ	Preston
700	335364	PR1 3YH	Preston
704	726619	PR1 5XA	Preston
719	479766	PR1 2EE	Preston
752	121893	PR2 6YS	Preston
787	92213	PR4 0BJ	Preston
821	136094	PR1 6EY	Preston
827	670391	PR1 6QY	Preston
828	78740	PR2 3XA	Preston
829	510354	PR1 2UP	Preston
856	479530	PR1 4TA	Preston
863	479853	PR1 5RY	Preston
868	479778	PR1 5TR	Preston
871	108298	PR1 3YH	Preston
896	136644	PR1 6QA	Preston
912	83211	PR1 5AS	Preston

925	335327	PR2	9UP	Preston
936	94388	PR1	2AR	Preston
954	479819	PR1	5XA	Preston

A `DataFrame` is an object in the `Pandas` library, but in addition we have the `Series` object, a collection of which makes up the `DataFrame`.

Many of the operations we can perform on a `Series` can also be performed on a `DataFrame`. It is a `Series` object which we will be using this week.

It is possible to perform an operation on each element in the `Series`, as well as call functions which require all of these such as `mean()`.

In [7]:

```
print(first_1000['RatingValue'].head(), '\n')
print(first_1000['RatingValue'].head() * 100, '\n')
print(first_1000['RatingValue'].head() * 23 > 100)
```

```
0    3.0
1    5.0
2    5.0
3    2.0
4    5.0
Name: RatingValue, dtype: float64
```

```
0    300.0
1    500.0
2    500.0
3    200.0
4    500.0
Name: RatingValue, dtype: float64
```

```
0    False
1     True
2     True
3    False
4     True
Name: RatingValue, dtype: bool
```

As well as being a part of a `DataFrame`, it is possible to create a `Series` from a list type object, for example see the code below:

In [8]:

```
s = pd.Series([8,6,2,7,9,6]) #pd가 임포트 pandas를 의미하는거임
    #시리즈 함수를 이용하여 다음과 같은 값(리스트)을 시리즈에 할당(넣어줌)
print(type(s))
print(s)
```

#출력값 : 시리즈 첫번째, 멤버 0는 8이다.

```
<class 'pandas.core.series.Series'>
0      8
1      6
2      2
3      7
4      9
5      6
dtype: int64
```

Create a Series object `rating_series` which contains the `RatingValue` column from the `first_1000` DataFrame object.

Then display descriptive statistics from that object (mean, median, mode etc). You can see the full list of available functions in [the documentation \(http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.html\)](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.html).

In [27]:

```
# YOUR CODE HERE
rating_series = pd.Series(first_1000['RatingValue'])
mean = rating_series.mean()
median = rating_series.median()
mode = rating_series.mode() #첫번째 최빈값은 0의 순서인 5.0이다.
    #만약 같은 최빈값이 있다면 1의 순서인 ?가 나오겠지.

print(mean)
print(median)
print(mode)
```

```
4.29613733906
5.0
0      5.0
dtype: float64
```

## Bokeh Charts

Bokeh uses Pandas data structures as the basis for its charts. We will now place the data structure we just generated onto various types of charts.

In Bokeh, there is a class **Plot** which is the basis for the visualisations we will consider in more detail next week. A **Chart** is a subclass of **Plot**, which is designed to allow the generation of common charts with the minimum amount of code.

At its simplest, these charts simply accept an object such as a `DataFrame` or `Series`, and will output a chart. For example, a simple bar chart looks as follows:



In [2]:

```
# To display the chart in the notebook, we need to run this function, otherwise call
output_notebook() #보케.io에서 만들어준경
k = first_1000['RatingValue'].head()
# Create a simple bar chart with made up data
bar_chart = Bar(pd.Series(k)) #세트 혹은 리스트 멤버를 시리즈로 사용함!

# Display the bar chart
show(bar_chart) #보케.io에서 만들어준경
```

(<http://bokeh.pydata.org>)

There are other charts which you can use such as a Histogram, Line graph, and scatter plot. View the Bokeh [user guide for charts](http://bokeh.pydata.org/en/latest/docs/reference/charts.html) (<http://bokeh.pydata.org/en/latest/docs/reference/charts.html>) to see the options available to customise the chart created above. In addition it is possible to customise the display.

## Randomness

A feature which NumPy supports is that of generating random numbers. This is important, for example, in generating a random sample from a population. The randomness, however, is not *truly* random, but rather pseudo-random, i.e., it will generate predictable values based on the initial **seed** that it accepts.

This feature means that if we know the seed, we can reproduce the results we wish to share provided that we have the same data, which is a desirable property. Though this may sound counter-intuitive it allows others to run our code using the same seed and they will get the same output, the code can then be run using a different seed value.

Consider the scenario where you want to populate an array with random data, you can use the `numpy.random.randint` function as below:

In [7]:

```
# import numpy.random
import numpy as np ##넘파이에서 랜덤 만들기
# The numbers generated will include the low value
low = 0
# The numbers generated will not include the high value, but will go up to high - 1
high = 10
np.random.randint(low, high, size=10)
#랜드int 평션 이용하기 random의 자식임, random은 np의 자식이고 !

print(np.random.randint(low, high, size=10))
```

[5 6 2 3 1 7 9 9 3 2]

Create a loop with 10 iterations, where each iteration prints a randomly generated array of size 10. Notice how each array has set of different values.

In [8]:

```
# YOUR CODE HERE
for ri in range(0,10):
    print(np.random.randint(low, high, size=10))
```

```
[1 1 8 4 9 6 8 4 6 8]
[3 0 6 4 2 8 3 0 4 7]
[8 3 8 2 4 5 3 6 2 4]
[9 0 9 7 4 9 4 9 0 7]
[4 5 2 9 2 2 0 8 8 1]
[2 6 8 2 3 4 2 0 0 7]
[9 7 4 5 9 6 5 2 6 9]
[3 7 1 7 4 4 3 9 2 7]
[8 8 9 2 2 0 3 9 6 1]
[4 5 6 4 8 9 8 7 5 7]
```

For many situations, this is desirable. However, where we want to be able to reproduce, e.g., sample sizes, we want our samples to be reproducible. To do this, we use the `RandomState` class in NumPy, where we specify our seed, as follows:

In [9]:

```
# Run this cell several times - observe the outcome
rs = np.random.RandomState(543210)
j = rs.randint(low, high, size=10) #씨드벨류 543210을 넣으면 사용자가 전달한 시드값을 기반으로 함
print(j)
```

```
[8 4 1 0 7 6 4 6 8 3]
```

In the cell below, generate the same loop as before, except this time instantiate the `RandomState` object to a value of 123456:

In [40]:

```
# YOUR CODE HERE
rs = np.random.RandomState(123456)
j = rs.randint(low, high, size=10)
print(j)
```

```
[1 2 1 8 0 7 4 8 4 2]
```

## IQR and Outliers

In the videos, Sergej talked about "outliers" in a dataset. In this worksheet, we'll give a slightly more detailed definition about what exactly they are, and the effect they can have on data.

An outlier is a value which is atypical of the rest of the dataset. For example, consider this set of data from [searches on the UK income tax calculator \(https://www.incometaxcalculator.org.uk/average-salary-uk.php\)](https://www.incometaxcalculator.org.uk/average-salary-uk.php). If we draw a distribution of them, we will notice a big difference in the values:

In [11]:

```

import pandas as pd
import numpy as np
from bokeh.charts import Histogram, output_notebook, show
from bokeh.models import Axis

salaries_list = [30000,18000,25000,20000,40000,50000,35000,45000,22000,60000,24000,
16000,100000,21000,26000,15000,32000,19000,17000,70000,27000,55000,18500,
36000,65000,42000,38000,12000,2481300,75000,33000,19500,43000,48000,12000,
17500,90000,34000,29000,16500,11000,31000,150000,37000,13000,22500,52000,
44000,200000,39000,46000,110000,27500,21500,47000,23500,15500,41000,26500,
20500,14500,130000,250000,24500,28500,72000,140000,32500,8000,53000,95000,
#50개의 결과(데이터)가 있는 검색 1번
print(salaries_list)
salaries = pd.DataFrame({'Salaries': salaries_list})

from bokeh.plotting import figure, show, output_file
output_notebook()
hist = Histogram(salaries, bins=50)
# Show absolute number on axis rather than E notation:
xaxis = hist.select(dict(type=Axis))[0] ##축 서식 지정에 관한 정보
xaxis.formatter.use_scientific = False ##축 서식 지정에 관한 정보
show(hist)


```

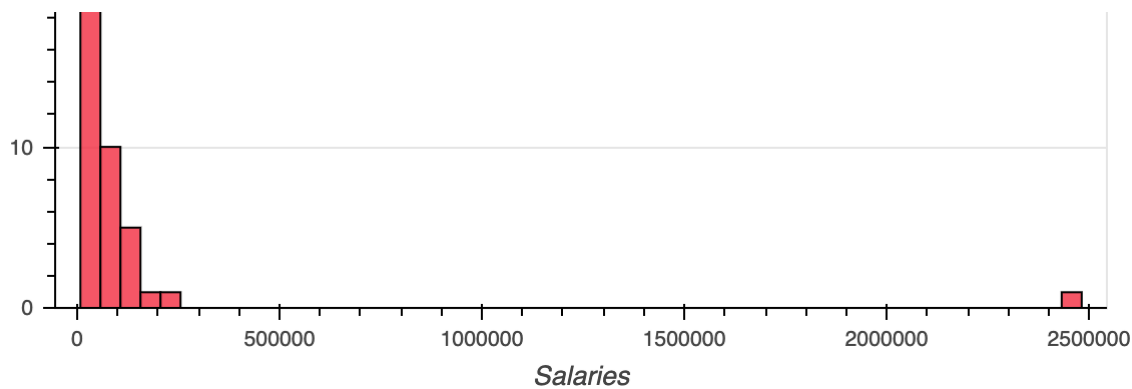
```

[30000, 18000, 25000, 20000, 40000, 50000, 35000, 45000, 22000, 60000,
24000, 28000, 23000, 16000, 100000, 21000, 26000, 15000, 32000, 19000,
17000, 70000, 27000, 55000, 18500, 80000, 36000, 65000, 42000, 38000,
12000, 2481300, 75000, 33000, 19500, 43000, 48000, 120000, 14000, 1750
0, 90000, 34000, 29000, 16500, 11000, 31000, 150000, 37000, 13000, 225
00, 52000, 10000, 85000, 44000, 200000, 39000, 46000, 110000, 27500, 2
1500, 47000, 23500, 15500, 41000, 26500, 15600, 16800, 20500, 14500, 1
30000, 250000, 24500, 28500, 72000, 140000, 32500, 8000, 53000, 95000,
25500]

```

(<http://bokeh.pydata.org/>) successfully loaded.





You will notice, the massive outlier on the right, where the person in question earns nearly £2.5 million. It makes it very difficult to get the chart to display anything useful, and has a significant effect on our data. For example, see the code below which shows the difference in number between the mean and the median:

In [12]:

```
print('The mean is: %f, and the median is %f' % (salaries.mean(), salaries.median()))
```

The mean is: 76371.250000, and the median is 31500.000000

Task: Remove the highest value from the dataset and see how this changes the mean and the median.

Further information about explicit location based indexing `.loc` can be found on this [Pandas Documentation](http://pandas.pydata.org/pandas-docs/version/0.15.0/indexing.html) (<http://pandas.pydata.org/pandas-docs/version/0.15.0/indexing.html>) page.

In [21]:

```
# YOUR CODE HERE 높은 값 제거
salaries = salaries.loc[salaries['Salaries'] < 2000000]
# loc[] 함수는 판다스의 위치기반 인덱싱 기능
print('The mean is: %f, and the median is %f' % (salaries.mean(), salaries.median()))
```

```
-----
-----
TypeError                                Traceback (most recent call
  last)
pandas/index.pyx in pandas.index.IndexEngine.get_loc (pandas/index.c:4
279)()
```

```
pandas/src/hashtable_class_helper.pxi in pandas.hashtable.Int64HashTab
le.get_item (pandas/hashtable.c:8543)()
```

TypeError: an integer is required

During handling of the above exception, another exception occurred:

```
KeyError                                Traceback (most recent call
  last)
<ipython-input-21-e60196e84ad2> in <module>()
      1 # YOUR CODE HERE 높은 값 제거
----> 2 salaries = salaries.loc[salaries['Salaries'] < 2000000]
      3 # loc[] 함수는 판다스의 위치기반 인덱싱 기능
      4 print('The mean is: %f, and the median is %f' % (salaries.mean
(), salaries.median()))
      5 print(salaries)
```

```
/opt/conda/lib/python3.5/site-packages/pandas/core/series.py in __geti
tem__(self, key)
    601         key = com._apply_if_callable(key, self)
    602         try:
--> 603             result = self.index.get_value(self, key)
    604
    605             if not is_scalar(result):
```

```
/opt/conda/lib/python3.5/site-packages/pandas/indexes/base.py in get_v
alue(self, series, key)
    2167         try:
    2168             return self._engine.get_value(s, k,
-> 2169                                     tz=getattr(series.dt
ype, 'tz', None))
    2170         except KeyError as e1:
    2171             if len(self) > 0 and self.inferred_type in ['integ
er', 'boolean']:
```

```
pandas/index.pyx in pandas.index.IndexEngine.get_value (pandas/index.
c:3557)()
```

```
pandas/index.pyx in pandas.index.IndexEngine.get_value (pandas/index.
c:3240)()
```

```
pandas/index.pyx in pandas.index.IndexEngine.get_loc (pandas/index.c:4
363)()
```

KeyError: 'Salaries'

Moving the top value had a considerable effect on the mean value of the dataset, decreasing it by over £30,000, however the result is still quite a bit higher than the median. So although the £2.5 million figure is obviously an outlier, how can we define an outlier more concretely? To start, we will consider the interquartile range (IQR):

## IQR

The IQR is calculated as follows:

1. Ordering the data by value
2. Taking the middle value from the *bottom* half of the data (lower quartile, known as Q1)
3. The median is known as Q2
4. Taking the middle value from the *top* half of the data (upper quartile, known as Q3)
5. The IQR is then calculated with  $Q3 - Q1$

The Q1 and Q2 values are considered as the 25th and 75th percentiles, since they represent the values 25% and 75% through the ordered data. Luckily, there are functions within Pandas which allow the calculation of these percentiles, which provide us with the IQR: the `quantile` (<http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.quantile.html>) function on a `DataFrame` which takes a float between 0 and 1 to get the appropriate percentile. For example, the median could be calculated as follows:

In [14]:

```
salaries.quantile(0.5) #중앙값, 중앙값 사분위수 31000
```

Out[14]:

```
Salaries      31000.0
Name: 0.5, dtype: float64
```

Task: Calculate the IQR of the salaries data using Pandas. Print out the value of the upper and lower quartiles so you can check the answer is correct

In [15]:

```
# YOUR CODE HERE
upper = float(salaries.quantile(0.75)) #상위 사분위수 51000
lower = float(salaries.quantile(0.25)) #하위 사분위수 20250
iqr = upper - lower #사분범위 IQR = 30750
print('Upper quartile:\t%f\nLower quartile:\t%f\nIQR:\t\t%f' % (upper, lower, iqr))

Upper quartile: 51000.000000
Lower quartile: 20250.000000
IQR:           30750.000000
```

## Outliers

Having introduced the IQR, we can now consider what constitutes an outlier. As a rule of thumb, an outlier can be defined as follows:

- $\text{lower\_quartile} - (1.5 * \text{IQR})$
- $\text{upper\_quartile} + (1.5 * \text{IQR})$

This is highly dependent on the data, and may not be appropriate for all situations, as is the decision of what to do with them. For the time being, we will simply exclude data which are outside these limits.

To do this, consider the following Pandas code, which excludes outliers from the salaries data. It uses a more complicated version of `.loc`, where it filters on two conditions

In [22]:

```
## 2000000 제거 하는 거 보다 더 좋은 방법 !

## lower_quartile - (1.5 * IQR), upper_quartile + (1.5 * IQR) 한계치 보다
## (상한 하한 보다 ) 높거나 낮은 모든 것은 이상치로 간주함, 항상 적용되는건 아님
#You don't need to write anything here
# Create the dataset again, rather than use the one with the top value taken out
salaries = pd.DataFrame({'Salaries': salaries_list})

salaries = salaries['Salaries'][(salaries['Salaries'] > (float(lower) - (iqr * 1.5))
                                & (salaries['Salaries'] < (float(upper) + (iqr * 1.5)))]

print('The mean is: %f, and the median is %f' % (salaries.mean(), salaries.median()))

print(salaries)
hist = Histogram(salaries, bins=50)
show(hist)
```

The mean is: 34202.816901, and the median is 28000.000000

0	30000
1	18000
2	25000
3	20000
4	40000
5	50000
6	35000
7	45000
8	22000
9	60000
10	24000
11	28000
12	23000
13	16000
15	21000
16	26000
17	15000
18	32000
19	19000
20	17000
21	70000
22	27000
23	55000
24	18500
25	80000
26	36000
27	65000
28	42000
29	38000
30	12000
	...
44	11000
45	31000
47	37000
48	13000
49	22500
50	52000
51	10000
52	85000
53	44000
55	39000




```

56 46000
58 27500
59 21500
60 47000
61 23500
62 15500
63 41000
64 26500
65 15600
66 16800
67 20500
68 14500
71 24500
72 28500
73 72000
75 32500
76 8000
77 53000
78 95000
79 25500

```

Name: Salaries, dtype: int64



The purpose of this exercise was to introduce the concept of an outlier, and how much of an effect it can have on data, and to give some practice using Pandas. There are many different ways that outliers could be defined,

and circumstances where they could or should not be excluded.